

**Please note this kernel is for practice purposes only.**

In this kernel I will be using AlexNet (<https://en.wikipedia.org/wiki/AlexNet>) for multiclass image classification.

**Inferences from the given dataset description:**

- There are 20,580 dogs images divided into 120 different categories (i.e., 120 breeds of dogs)

**Steps followed in this kernel:**

- Pick different categories of dog images for training the CNN model.
- After the training is done we will check the accuracy of the CNN model in predicting the dog's breed correctly given an image of the dog.

## 1. Load all the libraries

```
In [1]:
import os
import cv2
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

from tensorflow import keras
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
from keras.layers.normalization import BatchNormalization

print("Loaded all libraries")
```

Loaded all libraries

Using TensorFlow backend.

## 2. Data loading and exploration

```
In [2]:
fpath = "../input/images/Images/"
random_seed = 42

categories = os.listdir(fpath)
categories = categories[:20]
print("List of categories = ", categories, "\n\nNo. of categories = "
, len(categories))
```

```
List of categories = ['n02105162-malinois', 'n02094258-Norwich_terrier', 'n02102177-Welsh_springer_spaniel', 'n02086646-Blenheim_spaniel', 'n02086910-papillon', 'n02093256-Staffordshire_bullterrier']
```

```
'n02113624-toy_poodle', 'n02105056-groenendael', 'n02109961-Eskimo_
dog', 'n02116738-African_hunting_dog', 'n02096177-cairn', 'n0209658
5-Boston_bull', 'n02100735-English_setter', 'n02102973-Irish_water_
spaniel', 'n02099429-curly-coated_retriever', 'n02088364-beagle',
'n02101006-Gordon_setter', 'n02108089-boxer', 'n02097130-giant_schn
auzer', 'n02112137-chow']
```

No. of categories = 20

In [3]:

```
def load_images_and_labels(categories):
    img_lst=[]
    labels=[]
    for index, category in enumerate(categories):
        for image_name in os.listdir(fpath+"/"+category):
            img = cv2.imread(fpath+"/"+category+"/"+image_name)
            img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

            img_array = Image.fromarray(img, 'RGB')

            #resize image to 227 x 227 because the input image resolu
tion for AlexNet is 227 x 227
            resized_img = img_array.resize((227, 227))

            img_lst.append(np.array(resized_img))

            labels.append(index)
    return img_lst, labels

images, labels = load_images_and_labels(categories)
print("No. of images loaded = ",len(images),"No. of labels loaded
= ",len(labels))
print(type(images),type(labels))
```

```
No. of images loaded = 3337
No. of labels loaded = 3337
<class 'list'> <class 'list'>
```

In [4]:

```
images = np.array(images)
labels = np.array(labels)

print("Images shape = ",images.shape,"Labels shape = ",labels.sha
pe)
print(type(images),type(labels))
```

```
Images shape = (3337, 227, 227, 3)
Labels shape = (3337,)
<class 'numpy.ndarray'> <class 'numpy.ndarray'>
```

- Check few random images and labels by displaying them in a graph

In [5]:

```
def display_rand_images(images, labels):
    plt.figure(1, figsize = (19, 10))
    n = 0
```

```

for i in range(9):
    n += 1
    r = np.random.randint(0, images.shape[0], 1)



    plt.subplot(3, 3, n)
    plt.subplots_adjust(hspace = 0.3, wspace = 0.3)
    plt.imshow(images[r[0]])

    plt.title('Dog breed : {}'.format(labels[r[0]]))
    plt.xticks([])
    plt.yticks([])

plt.show()

display_rand_images(images, labels)

```



### 3. Prepare data for training the CNN model

- For training the CNN model we have to shuffle all the data that is loaded in images, labels list.

In [6]:

```

#1-step in data shuffling

#get equally spaced numbers in a given range
n = np.arange(images.shape[0])
print("'n' values before shuffling = ",n)

#shuffle all the equally spaced values in list 'n'
np.random.seed(random_seed)
np.random.shuffle(n)
print("\n'n' values after shuffling = ",n)

```

'n' values before shuffling = [ 0 1 2 ... 3334 3335 3336]

'n' values after shuffling = [ 321 3095 727 ... 1294 860 3174]

In [7]:

*#2-step in data shuffling**#shuffle images and corresponding labels data in both the lists*

images = images[n]

labels = labels[n]

```
print("Images shape after shuffling = ",images.shape,"\nLabels shape after shuffling = ",labels.shape)
```

Images shape after shuffling = (3337, 227, 227, 3)

Labels shape after shuffling = (3337,)

- Data normalization

In [8]:

images = images.astype(np.float32)

labels = labels.astype(np.int32)

images = images/255





```
print("Images shape after normalization = ",images.shape)
```

Images shape after normalization = (3337, 227, 227, 3)

- Display few random images after data normalization

In [9]:

display\_rand\_images(images, labels)



- Split the loaded dataset into train, test sets

In [10]:

```
x_train, x_test, y_train, y_test = train_test_split(images, labels,
test_size = 0.2, random_state = random_seed)
```

```
print("x_train shape = ",x_train.shape)
```

```
print("y_train shape = ",y_train.shape)
```

```
print("\nx_test shape = ",x_test.shape)
print("y_test shape = ",y_test.shape)
```

```
x_train shape = (2669, 227, 227, 3)
y_train shape = (2669,)
```

```
x_test shape = (668, 227, 227, 3)
y_test shape = (668,)
```

In [11]:

```
display_rand_images(x_train, y_train)
```



## 4. Define AlexNet CNN model

- Define all layers in the AlexNet CNN model

In [12]:

```
model=Sequential()

#1 conv layer
model.add(Conv2D(filters=96, kernel_size=(11,11), strides=(4,4), padding="valid", activation="relu", input_shape=(227,227,3)))

#1 max pool layer
model.add(MaxPooling2D(pool_size=(3,3), strides=(2,2)))

model.add(BatchNormalization())

#2 conv layer
model.add(Conv2D(filters=256, kernel_size=(5,5), strides=(1,1), padding="valid", activation="relu"))

#2 max pool layer
model.add(MaxPooling2D(pool_size=(3,3), strides=(2,2)))

model.add(BatchNormalization())
```

```

#3 conv layer
model.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding="valid", activation="relu"))

#4 conv layer
model.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding="valid", activation="relu"))

#5 conv layer
model.add(Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), padding="valid", activation="relu"))

#3 max pool layer
model.add(MaxPooling2D(pool_size=(3,3), strides=(2,2)))

model.add(BatchNormalization())

model.add(Flatten())

#1 dense layer
model.add(Dense(4096, input_shape=(227, 227, 3), activation="relu"))

model.add(Dropout(0.4))

model.add(BatchNormalization())

#2 dense layer
model.add(Dense(4096, activation="relu"))

model.add(Dropout(0.4))

model.add(BatchNormalization())

#3 dense layer
model.add(Dense(1000, activation="relu"))

model.add(Dropout(0.4))

model.add(BatchNormalization())

#output layer
model.add(Dense(20, activation="softmax"))

model.summary()

```

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/tensorflow/python/framework/op\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/keras/backend/tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `ra

```
te = 1 - keep_prob`.
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 55, 55, 96)	34944
max_pooling2d_1 (MaxPooling2D)	(None, 27, 27, 96)	0
batch_normalization_1 (Batch Normalization)	(None, 27, 27, 96)	384
conv2d_2 (Conv2D)	(None, 23, 23, 256)	614656
max_pooling2d_2 (MaxPooling2D)	(None, 11, 11, 256)	0
batch_normalization_2 (Batch Normalization)	(None, 11, 11, 256)	1024
conv2d_3 (Conv2D)	(None, 9, 9, 384)	885120
conv2d_4 (Conv2D)	(None, 7, 7, 384)	1327488
conv2d_5 (Conv2D)	(None, 5, 5, 256)	884992
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 256)	0
batch_normalization_3 (Batch Normalization)	(None, 2, 2, 256)	1024
flatten_1 (Flatten)	(None, 1024)	0
dense_1 (Dense)	(None, 4096)	4198400
dropout_1 (Dropout)	(None, 4096)	0
batch_normalization_4 (Batch Normalization)	(None, 4096)	16384
dense_2 (Dense)	(None, 4096)	16781312
dropout_2 (Dropout)	(None, 4096)	0
batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
dense_3 (Dense)	(None, 1000)	4097000
dropout_3 (Dropout)	(None, 1000)	0
batch_normalization_6 (Batch Normalization)	(None, 1000)	4000
dense_4 (Dense)	(None, 20)	20020
Total params: 28,883,132		
Trainable params: 28,863,532		
Non-trainable params: 19,600		

- Compile the CNN model

```
In [13]:
```

```
model.compile(optimizer="adam", loss="sparse_categorical_crossentropy", metrics=["accuracy"])
```

## 5. Train the model

- Fit the model using training data

In [14]:

```
%%time
model.fit(x_train, y_train, epochs=100)
```

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/tensorflow/python/ops/math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Epoch 1/100

2669/2669 [=====] - 8s 3ms/step - loss: 3.4869 - acc: 0.1034

Epoch 2/100

2669/2669 [=====] - 3s 1ms/step - loss: 3.1226 - acc: 0.1278

Epoch 3/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.9775 - acc: 0.1487

Epoch 4/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.8394 - acc: 0.1712

Epoch 5/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.7636 - acc: 0.1888

Epoch 6/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.7037 - acc: 0.2016

Epoch 7/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.6212 - acc: 0.2068

Epoch 8/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.5784 - acc: 0.2196

Epoch 9/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.4898 - acc: 0.2308

Epoch 10/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.4467 - acc: 0.2375

Epoch 11/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.3674 - acc: 0.2698

Epoch 12/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.3088 - acc: 0.2825

Epoch 13/100

2669/2669 [=====] - 3s 1ms/step - loss: 2.2494 - acc: 0.2900

Epoch 14/100



```
2669/2669 [=====] - 3s 1ms/step - loss: 2.
1742 - acc: 0.3080
Epoch 15/100
2669/2669 [=====] - 3s 1ms/step - loss: 2.
1548 - acc: 0.3245
Epoch 16/100
2669/2669 [=====] - 3s 1ms/step - loss: 2.
0757 - acc: 0.3421
Epoch 17/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
9860 - acc: 0.3720
Epoch 18/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
8994 - acc: 0.3810
Epoch 19/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
8457 - acc: 0.4061
Epoch 20/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
7737 - acc: 0.4215
Epoch 21/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
8074 - acc: 0.4125
Epoch 22/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
6133 - acc: 0.4792
Epoch 23/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
5159 - acc: 0.5051
Epoch 24/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
4338 - acc: 0.5257
Epoch 25/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
3129 - acc: 0.5643
Epoch 26/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
2373 - acc: 0.5995
Epoch 27/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
0869 - acc: 0.6373
Epoch 28/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
9776 - acc: 0.6767
Epoch 29/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
0310 - acc: 0.6557
Epoch 30/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
8480 - acc: 0.7123
Epoch 31/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
6950 - acc: 0.7628
Epoch 32/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
7164 - acc: 0.7636
Epoch 33/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
5980 - acc: 0.8048
Epoch 34/100
```

```
Epoch 34/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
5337 - acc: 0.8213
Epoch 35/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
4588 - acc: 0.8449
Epoch 36/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
4066 - acc: 0.8587
Epoch 37/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
4646 - acc: 0.8471
Epoch 38/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
3716 - acc: 0.8741
Epoch 39/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
3064 - acc: 0.9003
Epoch 40/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
3934 - acc: 0.8685
Epoch 41/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
9287 - acc: 0.4706
Epoch 42/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
9472 - acc: 0.6875
Epoch 43/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
9284 - acc: 0.6965
Epoch 44/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
8388 - acc: 0.7257
Epoch 45/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
4792 - acc: 0.8456
Epoch 46/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
3194 - acc: 0.9045
Epoch 47/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2686 - acc: 0.9090
Epoch 48/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2564 - acc: 0.9202
Epoch 49/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2178 - acc: 0.9281
Epoch 50/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2147 - acc: 0.9314
Epoch 51/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1998 - acc: 0.9314
Epoch 52/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1756 - acc: 0.9371
Epoch 53/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1945 - acc: 0.9329
```

```
Epoch 54/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2205 - acc: 0.9183
Epoch 55/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2696 - acc: 0.9030
Epoch 56/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2580 - acc: 0.9153
Epoch 57/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
7208 - acc: 0.8262
Epoch 58/100
2669/2669 [=====] - 3s 1ms/step - loss: 1.
0705 - acc: 0.6782
Epoch 59/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
3509 - acc: 0.8820
Epoch 60/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2033 - acc: 0.9363
Epoch 61/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1612 - acc: 0.9479
Epoch 62/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1332 - acc: 0.9569
Epoch 63/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1242 - acc: 0.9580
Epoch 64/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1117 - acc: 0.9659
Epoch 65/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1334 - acc: 0.9528
Epoch 66/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1038 - acc: 0.9652
Epoch 67/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1410 - acc: 0.9543
Epoch 68/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1381 - acc: 0.9558
Epoch 69/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2320 - acc: 0.9292
Epoch 70/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1810 - acc: 0.9404
Epoch 71/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1304 - acc: 0.9554
Epoch 72/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1587 - acc: 0.9505
Epoch 73/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1215 - acc: 0.9599
```

```
Epoch 74/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1347 - acc: 0.9520
Epoch 75/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1161 - acc: 0.9644
Epoch 76/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1489 - acc: 0.9513
Epoch 77/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1324 - acc: 0.9592
Epoch 78/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1210 - acc: 0.9580
Epoch 79/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
0931 - acc: 0.9682
Epoch 80/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1053 - acc: 0.9674
Epoch 81/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
0885 - acc: 0.9726
Epoch 82/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1057 - acc: 0.9618
Epoch 83/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2130 - acc: 0.9281
Epoch 84/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1264 - acc: 0.9599
Epoch 85/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1514 - acc: 0.9520
Epoch 86/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1286 - acc: 0.9573
Epoch 87/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
0922 - acc: 0.9700
Epoch 88/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
0664 - acc: 0.9771
Epoch 89/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
0865 - acc: 0.9741
Epoch 90/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
3682 - acc: 0.8973
Epoch 91/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
2416 - acc: 0.9217
Epoch 92/100
2669/2669 [=====] - 3s 1ms/step - loss: 0.
1733 - acc: 0.9475
Epoch 93/100
```